

HKUST-Department of Computer Science and Engineering
COMP3711: Design and Analysis of Algorithms
Final Examination - Fall 2018


Name:	Student ID	ITSC account:
-------	------------	---------------

- This is a closed book exam. It consists of 10 pages and 8 questions.
- Answer all the questions within the space provided on the examination paper.
- **ONLY WRITE ON THE FRONTS OF PAGES. SOLUTIONS WRITTEN ON THE BACKS OF PAGES WILL BE IGNORED.**
- No math sheet is required for this exam.

Prob#	1	2	3	4	5	6	7	8	Total
points									

1. Topological Order and Breadth/Depth First Traversal [10 points]

Consider the following directed acyclic graph.



Part (a) [4 points]

Provide **two** different valid topological orders.

Topological order 1: a, b, h, c, e, f, d


Topological order 2: a, b, c, h, e, f, d

Marking Scheme:

- i. 2 points for each correct topological order. There are some more correct topological orders: [a, b, h, e, c, f, d], [e, f, d, a, b, h, c], [e, f, d, a, b, c, h], ..., etc.
- ii. If topological order 1 is identical to topological order 2, count only once.
- iii. -1 point for any two out of order vertices for each topological order. E.g. if the answer of topological order 1 is [a, h, b, c, e, f, d], get 1 point out of 2.
0 point for more than 1 out of order pairs for each topological order.
- iv. If topological order 1 (or 2) is correct and topological order 2 (or 1) has exactly one out of order pair, and swapping the out of order vertices in topological order 2 (or 1) becomes identical to the topological order 1 (or 2), then the 1 point from (iii) will NOT be given because it's in fact two identical order.
 - Eg 1: topological order 1: [a, b, h, c, e, f, d] and topological order 2: [b, a, h, c, e, f, d]. Only 2 point for topological order 1, no point for topological order 2 because swapping (b, a) gives the same answer as topological order 1.
 - Eg 2: topological order 1: [a, b, h, c, e, f, d] and topological order 2: [e, d, f, a, b, h, c]. 2 point for topological order 1, 1 point for topological order 2 because topological order 2 has exactly one pair out of order (d, f). Swap (d, f) is different with topological order 1.

Part (b) [3 points]

Draw the BFS Tree if we start breadth first traversal from node a . (The tree only covers nodes reachable from a)




Marking Scheme:

-1 point for any extra or missing edge.

Part (c) [3 points]

Draw a DFS Tree if we start depth first traversal from node a . (The tree only covers nodes reachable from a)




Because the adjacency list is not given, there are more than one (total should be 4) possible depth first search results. Two possible solutions are given above.

Marking Scheme:

-1 point for any extra or missing edge.


2. Shortest Path and Minimum Spanning Trees [10 points]

Consider the following graph.



Part (a) [5 points]

Draw the *shortest path tree* generated by applying *Dijkstra's algorithm* starting from node *a*.




Marking Scheme:

- 0 point if not a tree that covers all vertices.
- 1 point for each wrong edge if the answer is a tree that covers all vertices.

Part (b) [5 points]

Draw the *minimum spanning tree* generated by applying *Prim's algorithm* starting from node *a*.



Marking Scheme:

- i. 0 point if not a tree that covers all vertices.
- ii. -1 point for each wrong edge if the answer is a tree that covers all vertices.

3. Greedy Algorithms [10 points]

Let $X = \{x_1, x_2, \dots, x_n\}$ be an input set of positive numbers such that $x_1 < x_2 < \dots < x_n$. Describe a $\Theta(n)$ algorithm that returns the minimum number of unit intervals (i.e., with length equal to 1) required to cover all numbers. We say that an interval $i = [i_s, i_e]$ covers x_i , if $i_s \leq x_i \leq i_e$.

Solution:

Keep the points in an array. Walk through the array as follows.

1. Set $x = x_1$
2. Walk through the points in increasing order
until finding the first j such that $x_j > x + 1$
3. Output
4. If there was no such j in (2) then stop. Otherwise, set $x = x_j$
5. Go to Step (2)

The running time is obviously $\Theta(n)$ because we just do a linear scan of the n numbers.

Marking Scheme:

- i. 0 points if the answer is $\omega(n \log n)$.
- ii. -7 points if the answer re-sort X (the given array is already sorted) or any correct $\Theta(n \log n)$ algorithm.
- iii. -3 points if missing “set $x = x_1$ ”.
- iv. -5 points if interval start at x_{j-1} instead of x_j for j defined in the solution above.
- v. -7 points if interval is not start at x_j for j defined in the solution above.
- vi. -1 if missing/incorrect running time analysis.
- vii. -1 point for each of any other careless mistake.
- viii. Prove of correctness is not necessary.

4. Stable Matching [12 points]

Consider the following preference list for four men and women.

Man	1 st Choice	2 nd Choice	3 rd Choice	4 th Choice
A	W	X	Y	Z
B	X	W	Z	Y
C	W	Z	X	Y
D	Z	X	W	Y

Woman	1 st Choice	2 nd Choice	3 rd Choice	4 th Choice
W	B	D	A	C
X	A	C	D	B
Y	D	B	C	A
Z	A	D	C	B

Part (a) [6 points]

Provide a stable matching, which is *optimal* for the men.

A-X, B-W, C-Y, D-Z

Part (b) [6 points]

Provide a stable matching, which is *optimal* for the women.

A-X, B-W, C-Y, D-Z

Marking Scheme for both (a) and (b):

-1.5 for each incorrect matching.

5. Divide and Conquer [14 points]

You are given an array $A[1..n]$ that contains a sequence of 0 followed by a sequence of 1 (e.g., 0001111111). A contains at least one 0 and one 1.

Part (a) [6 points]

Design an $O(\log n)$ algorithm that finds the position k of the last 0, i.e. $A[k] = 0$ and $A[k + 1] = 1$.

Solution:

find-k($A[p..r]$)

if $r = p + 1$, RETURN p

$mid = \lfloor (p+r)/2 \rfloor$

if $A[mid] = 0$ and $A[mid+1] = 1$, RETURN mid

if $A[mid] = 0$, **find-k**($A[mid..r]$)
 else **find-k**($A[p..mid]$)

This is similar to binary search: with a constant number of comparisons, we reduce the problem size by half: $T(n) = T(n/2) + c \Rightarrow T(n) = O(\log n)$

Marking Scheme:

- i. At most 1 point will be given if using Selection, Partition or Sorting.
- ii. -3 points for missing base case of both “if $r = p + 1$, RETURN p ” and “if $A[mid] = 0$ and $A[mid+1] = 1$, RETURN mid ”.
- iii. -2 points for missing base case of “if $r = p + 1$, RETURN p ”.
 - Note that, missing “if $A[mid] = 0$ and $A[mid+1] = 1$, RETURN mid ” is fine as long as “if $r = p + 1$, RETURN p ” is in the algorithm.
- iv. If missing “if $A[mid] = 0$ and $A[mid+1] = 1$, RETURN mid ”,
 - Some answer’s base case might be “if $r == p$ then, return p ”. In this case, the recursive call should be something like “If $A[mid] == 0$, **find-k**($A[mid..r]$) else **find-k**($A[p..mid-1]$)”
 Otherwise, the algorithm may never be terminated.
- v. -1 point for missing a brief explanation of running time.
- vi. -1 point for each of any other careless mistakes.
- vii. Prove of correctness is not necessary.

Part (b) [8 points]

Suppose that k is much smaller than n . Design an $O(\log k)$ algorithm that finds the position k of the last 0.

Hint: You may re-use the solution of Part (a) in the solution of Part (b).

Solution:

$i \leftarrow 1$

while $A[i] = 0$

$i \leftarrow 2i$

find-k($A[i/2+1..i]$)

The while loop will stop when it finds a 1. Since each time we double the value of i , the while loop performs $\log k$ iterations. The first 1 occurs somewhere between the positions $A[i/2+1]$ and $A[i]$. To find it, we call **find-k**($A[i/2+1..i]$), which has cost $\log(k/2) = O(\log k)$. Therefore, the total cost is $O(\log k)$.

Marking Scheme:

It's fine to assume there exist a $O(\log n)$ algorithm to solve Part (a) and use here even though the answer of Part (a) is wrong.

- i. 5 points for repeatedly multiply i by 2 until $A[i] = 1$.
- ii. 3 points for running time analysis.
- iii. -1 point for each of any other careless mistakes.
- iv. Prove of correctness is not necessary.

6. Maximum Flow [16 pts]

A project is partitioned into m tasks t_1, \dots, t_m . There are n candidate workers c_1, \dots, c_n and n lists l_1, \dots, l_n . List l_j contains the tasks that candidate c_j can work for. For example, $l_1 = \{t_3, t_5\}$ means candidate c_1 can only be assigned to t_3 or t_5 . You can assign at most a_i candidates to task t_i for $1 \leq i \leq m$. Each candidate can be assigned to at most 1 task. Design an algorithm for finding a candidate-task assignment that maximizes the number of working candidates under the limitations of a_i and l_j for $1 \leq i \leq m$ and $1 \leq j \leq n$. Note that some tasks may have no candidate.

Solution:

We first form a flow network G as follows:

- The m tasks form m vertices in the flow network, denote these m vertices by t_1, \dots, t_m . The n candidates form n vertices in the flow network, denote these n vertices by c_1, \dots, c_n .
- According to the list l_j , add edges from c_j to the vertices of the tasks in l_j with capacity 1.
- Add a source vertex s and a sink vertex t . Connect s to t_i with capacity a_i for all $1 \leq i \leq m$. Connect t to c_j for all $1 \leq j \leq n$ with capacity 1.

Apply Ford-Fulkerson algorithm with Edmonds Karp suggestion to find the maximum flow of G . The maximum flow value is the maximum number of candidates. The candidates who received a flow is assigned to the project. Trace the flow pushed to the vertex c_i , you know which task is assigned to candidate i .

The running time to construct the flow network is $O(n + m + nm)$ time because there are $O(n + m)$ vertices and $O(nm)$ edges. The time for computing maximum flow is $O((n + m)(nm)^2)$.

Marking Scheme:

- i. Any correct algorithm runs in $O((n + m)(nm)^2)$ is fine, not restricted to Max Flow.

If the answer is converted to max flow problem:

- ii. -8 points if everything correct but without mention the edge capacities.
- iii. If some edge capacities are mentioned:
 - -2 points if wrong/missing capacity is assigned for edges from candidates to sink.

- -3 points if wrong/missing capacity is assigned for edges from source to tasks.
- -3 points if wrong/missing capacity is assigned for edges from tasks to candidates.
- iv. -2 if not single source single sink.
- v. -3 if incorrectly connect to source and/or sink.
- vi. -3 if incorrectly connect candidates and tasks.
- vii. Prove of correctness and running time analysis are not necessary.

7. Dynamic Programming [18 pts]

Consider a version of the 0-1 knapsack problem, where there are infinite copies of each item. Specifically, the input is a set of n items, where item i has weight w_i and value v_i , and a knapsack with capacity W . The goal is to find $x_1, \dots, x_n \in \{0, 1, 2, 3, \dots\}$ such that $\sum_{i=1}^n x_i w_i \leq W$ and $\sum_{i=1}^n x_i v_i$ is maximized. The difference with respect to the conventional 0-1 knapsack is that we can carry as many copies of the same item as we wish.

Part (a) [6 points]

Describe a greedy algorithm for solving the problem, and discuss if it is optimal or not.

Solution:

Take as many copies of the item with the largest $\frac{v_i}{w_i}$ as you can. Then, continue with the next most precious item that fits in the knapsack.

This will not give the optimal solution. Assume two items with $v_1 = 10$, $w_1 = 10$ and $v_2 = 3.5$, $w_2 = 4$. $W = 12$. The greedy solution will take item 1 with total value 10. The optimal solution will take 3 copies of item 2, with value 10.5.

Marking Scheme:

- i. 0 point for non-greedy solution
 - A greedy solution is not necessary to use $\frac{v_i}{w_i}$ as the key. It's fine to pick largest v_i or smallest w_i , etc, as long as it is always pick the largest or smallest element.
- ii. -4 points if saying greedy algorithm can yield optimal solution.
- iii. -2 points if saying greedy algorithm cannot yield optimal solution, but missing or incorrect example.
- iv. Proof of correctness and running time analysis are not necessary.

Part (b) [6 points]

Describe and explain the recurrence for a dynamic programming algorithm that finds the optimal solution.

Solution:

If Optimal Solution for knapsack of size j chooses item i , remainder of optimal solution is optimal solution for subproblem of filling knapsack of size $j - w_i$.

Recursive case: $V[j] = \max(0, v_1 + V[j - w_1], v_2 + V[j - w_2], \dots, v_n + V[j - w_n])$

Base cases: $V[j] = 0, j = 0$

$V[j] = -\infty, j < 0$

Marking Scheme:

- i. 0 point if EACH $V[j]$ is solved in exponential time (Note, $O(W)$ is also treated as exponential).

- ii. 1 point if completely correct and EACH $V[j]$ is solve in time polynomial of n , but $\omega(n)$.
- iii. -2 points for each missing/incorrect base case.
- iv. -2 points for incorrect recursive case.
 - Only -0.5 points if everything correct but missing 0, i.e.
$$V[j] = \max(v_1 + V[j - w_1], v_2 + V[j - w_2], \dots, v_n + V[j - w_n])$$
- v. Proof of correctness and running time analysis are not necessary.

Part (c) [6 points]

Write the pseudo-code based on the recurrence of Part (b).

Solution:

```
V[j] = 0
for j = 1 to W
    V[j] = 0
    for i = 1 to n
        if j ≥ wi and vi + V[j - wi] ≥ V[j]
            V[j] = vi + V[j - wi]
```

Marking Scheme:

- i. 0 point if completely not following the student's answer in Part (b).
- ii. 0 point if Part (b) is completely not solving the problem.
- iii. If Part (b) solves the problem (suffer minor mistake)
 - Full marks for pseudo-code exactly follows the student's answer in Part (b).
 - -1 point for any careless mistake.

8. Dynamic Programming over Intervals [10 points]

Given an array A of n non-negative integers and an integer m (m is much smaller than n), we wish to split A into m **non-empty continuous** subarrays, so that the largest sum among these m subarrays is minimized.

Example:

Assume that the input array is $A=[7,2,5,10,8]$ and $m=2$. There are four ways to split A into two subarrays: **(1)** $[7],[2,5,10,8]$, **(2)** $[7,2],[5,10,8]$, **(3)** $[7,2,5],[10,8]$, **(4)** $[7,2,5,10],[8]$. The best way is to split it into $[7,2,5]$ and $[10,8]$ because the largest sum between the two subarrays is only 18.

Provide and explain the recurrence for the optimal solution. You do **not** need to include any pseudo-code.

Solution:

Let $s[j] = \sum_{k=1}^j A[k]$, $s[j]$ can be computed in $O(n)$ time by

$$s[0] = 0, s[j] = s[j-1] + A[j]$$

With this pre-sum array s , $\sum_{k=i}^j A[k] = s[j] - s[i-1]$ can be computed in $O(1)$ time for $1 \leq i \leq j \leq n$.

For $1 \leq i \leq m$, $1 \leq j \leq n$, define $d[i][j]$ be the optimal value of the subproblem that we split $A[1..j]$ into i non-empty continuous subarrays so that the largest sum among these i subarrays is minimized.

Base case: $d[1][1] = A[1]$.

Recursive case: $d[i][j] = \min_{i-1 \leq k < j} (\max(d[i-1][k], s[j] - s[k]))$

The last split can be at position $i-1$ to $j-1$. When the last split is at position k , this means we have split the array into $i-1$ subarrays for the first k element, $d[i-1][k]$ stores the largest sum among these $i-1$ subarrays. The sum of the i -th subarray is $s[j] - s[k]$. So, $\max(d[i-1][k], s[j] - s[k])$ is the largest sum among the i subarrays subject to the last split of $A[1..j]$ is at position k . We simply try all possibility (all possible k) and pick the minimum.

The running time of this algorithm is $O(n^2m)$ because there are $O(nm)$ subproblems and each subproblem can be computed in $O(n)$ time. The time for computing s at the beginning is just $O(n)$.

Marking Scheme:

- i. 0 point for any answer that runs in $\Omega(n^4m)$.
- ii. -2 points for missing the pre-sum array s . Because it results in a very slow algorithm.
- iii. -2 points for missing/incorrect base case.
 - $d[0][0] = 0$ is a correct base case.
- iv. -3 points for $d[i][j] = \min_{i-1 \leq k < j} (d[i-1][k])$.
- v. -3 points for $d[i][j] = \min_{i-1 \leq k < j} (s[j] - s[k])$.
- vi. -1 for each careless mistake. E.g. $\min_{i \leq k < j}$ or $\min_{i-1 \leq k \leq j}$, etc.
- vii. Proof of correctness and running time analysis are not necessary.